

# Master in Data Science

Language  
Detection

General  
Structure

Detailed  
Structure

Core task

Goals &  
Deliverables

## Mining Unstructured Data

# Outline

Language  
Detection

General  
Structure

Detailed  
Structure

Core task

Goals &  
Deliverables

## 1 Language Detection

## 2 General Structure

## 3 Detailed Structure

## 4 Core task

## 5 Goals & Deliverables

# Session 1 - Language Detection

## Assignment

trying to predict what language the text is written is.

Study the impact of different preprocessing techniques on NLP task. To do so, we will perform **Language Detection** over the provided csv file *data.csv* and detects in language each sentence is written.

	Text	Language
0	klement gottwaldi surnukeha palsameeriti ning ...	Estonian
1	sebes joseph pereira thomas pâ eng the jesuit...	Swedish
2	ถนนเจริญกรุง อักษรโรมัน thanon charoen krung ...	Thai
3	விசாகப்பட்டினம் தமிழ்ச்சங்கத்தை இந்துப் பத்திர...	Tamil
4	de spons behoort tot het geslacht haliclona en...	Dutch
...	...	...
21995	hors du terrain les années et sont des année...	French
21996	ใน พศ หลักจากที่เสด็จประพาสแหลมมลายู ชาว ฮิน...	Thai
21997	con motivo de la celebración del septuagésimoq...	Spanish
21998	年月，當時還只有歲的她在美國出道，以mai-k名義推出首張英文《baby i like》，由...	Chinese
21999	aprilie sonda spațială messenger a nasa și-a ...	Romanian

22000 rows × 2 columns

Language  
Detection

General  
Structure

Detailed  
Structure

Core task

Goals &  
Deliverables

# Outline

Language  
Detection

General  
Structure

Detailed  
Structure

Core task

Goals &  
Deliverables

1 Language Detection

2 General Structure

3 Detailed Structure

4 Core task

5 Goals & Deliverables

# General Structure - Main function

Language  
Detection

General  
Structure

Detailed  
Structure

Core task

Goals &  
Deliverables

The program (langdetect.py) expects three arguments: Path to the input data, vocabulary size and analyzer granularity (words or chars)

```
def get_parser():  
    parser = argparse.ArgumentParser()  
    parser.add_argument("-i", "--input",  
                        help="Input data in csv format", type=str)  
    parser.add_argument("-v", "--voc_size",  
                        help="Vocabulary size", type=int)  
    parser.add_argument("-a", "--analyzer",  
                        help="Tokenization level: {word, char}",  
                        type=str, choices=['word', 'char'])  
    return parser
```

vocabulary size: how many words the system understands and how many resources are being used by the system. It influences the size of the vector.

# General Structure - Main function II

Then, it reads all files in the given directory, and splits the data in train a test splits

```
raw = pd.read_csv(args.input)

# Languages
languages = set(raw['language'])
print('=====')
```

There are 22 languages in our set

```
print('Languages', languages)
print('=====')
```

  

```
# Split Train and Test sets
X=raw['Text']
y=raw['language']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
                                                    random_state=seed)

print('=====')
```

```
print('Split sizes:')
print('Train:', len(X_train))
print('Test:', len(X_test))
print('=====')
```

Language  
Detection

General  
Structure

Detailed  
Structure

Core task

Goals &  
Deliverables

# General Structure - Main function III

Language  
Detection

General  
Structure

Detailed  
Structure

Core task

Goals &  
Deliverables

Then, it preprocesses the data, computes its features and the coverage of the vocabulary over the test data

```
# Preprocess text (Word granularity only)    add preprocessing pipeline here
    if args.analyzer == 'word':
        X_train, y_train = preprocess(X_train,y_train)
        X_test, y_test = preprocess(X_test,y_test)

#Compute text features
features, X_train_raw, X_test_raw = compute_features(X_train, X_test,
    analyzer=args.analyzer,max_features=args.voc_size)

print('=====' )
print('Number of tokens in the vocabulary:', len(features))
print('Coverage: ', compute_coverage(features, X_test.values, analyzer=
    args.analyzer))
print('=====' )
```

# General Structure - Main function IV

Language  
Detection

General  
Structure

Detailed  
Structure

Core task

Goals &  
Deliverables

Finally, it train a classifier model, predicts over the test set, reports its performance and plots its PCA dimensionality reduction

```
#Apply Classifier
```

```
X_train, X_test = normalizeData(X_train_raw, X_test_raw)
```

```
y_predict = applyNaiveBayes(X_train, y_train, X_test)
```

```
print('=====')
```

```
print('Prediction Results:')
```

```
plot_F_Scores(y_test, y_predict)
```

```
print('=====')
```

try different classifiers, but spend  
more efforts on the preprocessing  
pipeline

```
plot_Confusion_Matrix(y_test, y_predict, "Greens")
```

```
#Plot PCA
```

```
print('=====')
```

```
print('PCA and Explained Variance:')
```

```
plotPCA(X_train, X_test, y_test, languages)
```

```
print('=====')
```



# Outline

Language  
Detection

General  
Structure

Detailed  
Structure

Core task

Goals &  
Deliverables

1 Language Detection

2 General Structure

3 Detailed Structure

4 Core task

5 Goals & Deliverables

# Functions - Tokenize text

This function is currently empty, you can apply all preprocessing steps. **Resources you may use: NLTK, Spacy**

```
#Tokenizer function. You can add here different
preprocesses.
def preprocess(sentence, labels):
    '''
    Task: Given a sentence apply all the required
    preprocessing steps
    to compute train our classifier, such as sentence
    splitting,
    tokenization or lemmatization.

    Input: Sentence in string format
    Output: Preprocessed sentence either as a list or a
    string
    '''
    # Place your code here
    # Keep in mind that sentence splitting affects the
    number of sentences
    # and therefore, you should replicate labels to match
    .
    return sentence, labels
```

Language  
Detection

General  
Structure

Detailed  
Structure

Core task

Goals &  
Deliverables

# Functions - Classifier models

```
# You may add more classifier methods replicating this
function
def applyNaiveBayes(X_train, y_train, X_test):
    '''
    Task: Given some features train a Naive Bayes
    classifier
           and return its predictions over a test set
    Input; X_train -> Train features
           y_train -> Train_labels
           X_test  -> Test features
    Output: y_predict -> Predictions over the test set
    '''

    trainArray = toNumpyArray(X_train)
    testArray = toNumpyArray(X_test)

    clf = MultinomialNB()
    clf.fit(trainArray, y_train)
    y_predict = clf.predict(testArray)
    return y_predict
```

Language  
Detection

General  
Structure

Detailed  
Structure

Core task

Goals &  
Deliverables

# Outline

Language  
Detection

General  
Structure

Detailed  
Structure

Core task

Goals &  
Deliverables

1 Language Detection

2 General Structure

3 Detailed Structure

4 Core task

5 Goals & Deliverables

# Language Detection - First baseline

Compare two different configurations of the system.

Without modifying the code run the following configurations and compare their vocabulary coverage and performance. Explain why they show different error patterns.

- **Character level:**

vocabulary 1000

```
python langdetect.py -i dataset.csv -v 1000 -a char
```

- **Word level:**

```
python langdetect.py -i dataset.csv -v 1000 -a word
```

Language  
Detection

General  
Structure

Detailed  
Structure

Core task

Goals &  
Deliverables

# 1st Exercise - First Baseline

Language  
Detection

General  
Structure

Detailed  
Structure

Core task

Goals &  
Deliverables

Focus on the following elements to explain the behavior:

- How well does the vocabulary covers the data?
- Which languages produce more errors? What do they have in common (family, script, etc)?
- How languages overlap on the PCA plot? What could that overlapping mean?

## 2nd Exercise - Document Structure

Language  
Detection

General  
Structure

Detailed  
Structure

Core task

Goals &  
Deliverables

### Second exercise

- Try different vocabulary sizes and preprocessing steps to analyze the behavior of this kind of data.
- Improving F1 score is **NOT** the objective of the task. Focus on understanding how different parameters affect the results.

Understanding the results of different actions is the purpose of this task.

# Outline

Language  
Detection

General  
Structure

Detailed  
Structure

Core task

Goals &  
Deliverables

1 Language Detection

2 General Structure

3 Detailed Structure

4 Core task

5 Goals & Deliverables



# Deliverables

Language  
Detection

General  
Structure

Detailed  
Structure

Core task

Goals &  
Deliverables

Write a report describing the work carried out in this exercise.

The report must be a **single self-contained PDF document**, under 10 pages, containing:

- *Introduction*: What is this report about. What is the goal of the presented work.
- *Preprocess*: Describe the preprocessing steps tried and the rationale to employ them.
- *Code*: Include your preprocessing functions as well as classifiers in the document **Do not include any other code.** `preprocess.py` & `classifier.py`
- *Experiments and results*: Results obtained on the **test** datasets, for different rule combinations you deem relevant.  
**Keep result tables in the format produced by the program. You can just download them and use them in your document.**
- *Conclusions*: Final remarks and insights gained in this task.